

Unified Enterprise Knowledge Hub System

A Comprehensive Design for Document Modeling, Domain Schema,
Retrieval-Augmented LLM/Graph Architecture,
RASA Integration, and Automated API Learning,
and Enterprise Text Processing

Srikari Rallabandi

(Merged, Extended, and Enhanced from Prior Architecture Documents)

Monday 3rd March, 2025

Contents

1	Introduction	3
1.1	Key Motivations and Enhancements	3
2	Objectives and Scope	4
2.1	Objectives	4
2.2	Scope and Limitations	4
3	Architecture Overview	4
3.1	High-Level Diagram	5
3.2	Microservices-based NLP Library	5
4	Hierarchical Document Model for Scraped Data	6
4.1	Rationalization	6
4.2	Core Classes and Code Snippets	6
5	Domain Schema and JSON-Based Modeling	7
5.1	Rationalization	7
5.2	Schema Example	8
5.2.1	Parsing and Validation	8
6	Modular NLP Framework with Pluggable Components	9
6.1	Document Ingestion & Parsing	9
6.2	NLP Modules	9
6.3	Search Indexing & Storage	9
6.4	Query Interface & Applications	9
7	Linking Document Content and Domain Schema	11
7.1	Motivation	11
7.2	Process Flow	11
7.3	Example Linking Code	11

8	Embeddings and Retrieval-Augmented Generation (RAG)	11
8.1	Paragraph-Level vs. Domain Concept Embeddings	11
8.2	Implementation Steps	11
8.3	RAG Query Pipeline	11
8.4	Embeddings and RAG Diagram	12
9	Detailed “How” Steps & Best Practices	12
9.1	Ingestion and Parsing	12
9.2	Domain Schema Loading	12
9.3	Entity & Action Linking	12
9.4	Embedding Storage	12
9.5	RAG Query Processing	12
9.6	High-level Flow Diagram for How Steps	12
10	LLM Integration with Adapter Pattern	13
10.1	Adapter Pattern Design	13
10.2	Code Snippet: LLM Adapter Interface (Pseudocode)	13
11	Logical Components	14
11.1	Overview	14
11.2	Service Interaction Diagram	14
12	Real-Time and Batch Processing	14
12.1	Real-Time Processing	14
12.2	Batch Processing	14
13	RASA Integration and Automated API Learning	15
13.1	Why RASA?	15
13.2	API Learning: Generating RASA Artifacts	15
13.3	Sample Code for Dynamic Action Generation	15
13.4	Connecting Domain Schema to RASA	15
14	LLM Integration in the Framework	16
15	Additional Considerations	16
15.1	Evaluation and Benchmarking	16
15.2	Continuous Learning and Adaptation	16
15.3	Explainability and Human-in-the-Loop	17
16	Deployment, Scalability, and Testing	17
16.1	Error Handling and Monitoring	17
16.2	Scaling the Graph and Vector DB	17
16.3	Scaling LLM Inference	17
16.4	Testing and Validation	17
17	Security and Access Control	17
18	Conclusion, Critical Gaps, and Future Directions	17
18.1	Critical Gaps and Recommendations	18
18.2	Future Directions	18
19	Final Reflection	18

1 Introduction

This document merges and extends multiple detailed architectures into a single, **end-to-end Knowledge Hub design** for handling both **structured and unstructured data** at enterprise scale. It covers:

- A *hierarchical Document Model* for scraped or uploaded files (PDF, Word, ePub, HTML, images, video transcripts, etc.).
- A *Domain Schema* in JSON/YAML to drive entity linking and RASA integration.
- Advanced mechanisms for **embedding, relationship extraction, and retrieval-augmented generation (RAG)**.
- A **Knowledge Graph** to support semantic queries.
- **RASA-based conversational flows** and automated API learning.
- A **modular, pluggable library design for NLP tasks** as independent microservices.
- A robust LLM integration based on an Adapter Pattern (defaulting to Llama 3.3) for static model selection.
- Support for both real-time streaming (via Kafka/RabbitMQ) and batch processing (ETL).
- Built-in evaluation, benchmarking, and self-learning feedback loops.

This unified architecture aims to enable enterprise users to query unstructured text, access structured APIs, and receive context-aware, grounded responses.

By merging **Document structures**, a **Domain Schema**, **Graph-based data**, and a synergy of **RASA + LLM**, enterprise users can query unstructured text (via paragraphs/tables), access structured APIs, confirm context with domain-defined concepts, and produce grounded, multi-faceted responses.

1.1 Key Motivations and Enhancements

- **Unified Document Representation:** Ensure the preservation of the document's hierarchical structure, including the core components such as **Document**, **Chapter**, **Section**, and **Paragraph**, to maintain clarity and organization.
- **Domain Concept Integration:** Integrate domain-specific knowledge, such as **entities**, **actions**, and **interactions**, with extracted text references. This enhances retrieval capabilities and enables more advanced analytics on the document content.
- **RASA + LLM Synergy:** Leverage the strengths of both RASA's structured conversation flows and rule-based dialogues, alongside the flexibility of LLMs for open-ended Q&A and generative tasks. This synergy allows for more robust and dynamic conversational experiences.
- **API Learning:** Enable dynamic parsing and conversion of enterprise API definitions (e.g., OpenAPI, RAML, WSDL) into RASA actions and entities. This supports rapid integration of external data sources and APIs into the system.
- **Graph + Vector DB Synergy:** Combine the strengths of relationship-based queries in graph databases with semantic search capabilities through vector stores. This hybrid approach enhances data retrieval and relationship discovery.
- **RAG Query Flow:** Implement top-K embeddings for generating LLM responses that are grounded in real documents. This method mitigates the risks of hallucinations by ensuring factual accuracy in the output.
- **Scalability and Extensibility:** Design the system to be easily scalable, allowing for incremental updates to domain schemas, re-embedding of new text, and the integration of new APIs into RASA. The system also supports scaling of graph and vector retrieval as needed.

- **Modular NLP Microservices:** Provide independently deployable Python services for various NLP tasks, such as parsing, summarization, entity extraction, and more. This modularity supports flexibility and adaptability in system deployment.
- **Pluggable LLM Integration:** Implement an Adapter Pattern that allows enterprises to choose between different LLMs (with the default being Llama 3.3). This provides flexibility in model selection while maintaining a static model choice for consistency.
- **Scalable Real-Time and Batch Processing:** Support both event-driven and scheduled processing workflows to accommodate a variety of real-time and batch data processing needs.
- **Robust Evaluation and Self-Learning:** Employ industry-standard metrics (such as BLEU, F1, ROUGE, etc.) to evaluate system performance. Continuous improvement is facilitated through feedback loops, ensuring ongoing optimization of the system.

2 Objectives and Scope

2.1 Objectives

- **Modular Data Ingestion:** Handle diverse file types (PDF, Word, HTML, ePub, images, voice transcripts) plus domain-specific JSON/YAML schemas.
- **Hierarchical & Graph Storage:** Build an enterprise knowledge graph containing both *paragraph text* and domain *entities/relationships*.
- **RASA + LLM Integration:** Support structured conversation flows (API calls, data lookups) and open-domain chat & retrieval-augmented generation.
- **API Learning Automation:** Dynamically convert API specifications into RASA artifacts (intents, entities, custom actions).
- **NLP Microservices Integration:** Support both structured (RASA) and open-ended (LLM-based RAG) interactions.
- **API Learning Automation:** Dynamically convert API definitions into RASA artifacts.
- **Real-Time and Batch Processing:** Enable both immediate feedback (e.g. social media moderation) and end-of-day batch processing.
- **Self-Learning and Evaluation:** Continuously monitor and improve model performance with minimal human intervention.
- **Advanced Analytics (Optional):** Integrate temporal or causal reasoning, simulations, or GNN-based inferences as needed.

2.2 Scope and Limitations

- **Domain Schema Definition:** Must be maintained or updated by domain experts for accuracy.
- **Graph DB Scale:** Suitable for up to millions of nodes/relationships; beyond that, specialized partitioning or distributed topologies may be required.
- **LLM Context Window Constraints:** Large retrieval sets must be summarized or chunked prior to final LLM ingestion.
- **Optional GNN or Bayesian Modules:** Implemented if deeper link-prediction or causal analyses are needed (healthcare, supply chain simulations, etc.).

3 Architecture Overview

A top-level view of the integrated architecture is shown in Figure ?? . It merges the hierarchical document model, domain schema, knowledge graph, vector embeddings, and an RASA+LLM synergy.

3.1 High-Level Diagram

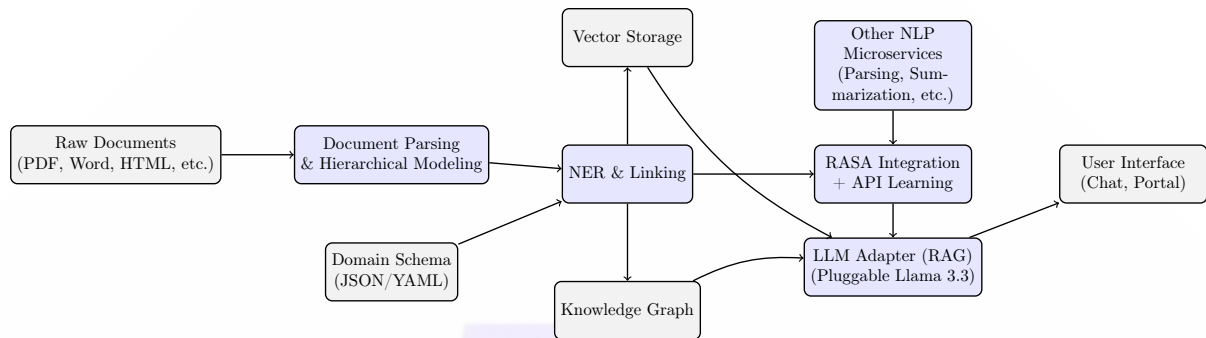


Figure 1: Enhanced high-level architecture: Document ingestion and domain schema processing feed a suite of independent NLP microservices. RASA handles structured interactions while the LLM Adapter (default Llama 3.3) enables retrieval-augmented generation.

The architecture comprises the following layers:

1. Data Ingestion and Parsing:

- Scrapers for structured data (APIs, DBs) or unstructured docs (PDF, DOCX, ePub, websites).
- Parsers produce **Document** objects with sub-structures (**Chapters**, **Sections**, **Paragraphs**, **Table**, **Image**).

2. Domain Schema Loading:

- JSON/YAML describing domain **entities**, **actions**, **interactions**.

3. NER and Linking:

- NLP-based NER to match textual references to domain **entities** and **actions**.

4. Embedding & Vector Store:

- Each **Paragraph** (or snippet) is converted to an embedding and stored in a vector database.

5. Graph Database:

- Store relationships (e.g., “Paragraph mentions Entity A”) for advanced queries.

6. RASA Integration & API Learning:

- Convert domain schema knowledge and API definitions into RASA artifacts.

7. RAG Query Flow:

- Convert user questions to embeddings, retrieve top-K paragraphs, and feed into an LLM.

3.2 Microservices-based NLP Library

The core of the enhanced architecture is a set of Python-based microservices that can be independently deployed and scaled. Key modules include:

- **Document Parsing Service:** Converts raw documents into hierarchical structures.
- **Image & Table Extraction Service:** Extracts and processes non-textual content.
- **Search and Indexing Service:** Builds traditional and vector-based search indexes.
- **Summarization Service:** Uses LLM-powered or rule-based summarization.
- **Context Awareness and Intent Recognition:** For chat and Q&A interfaces.

- **Sentiment and Spam Detection Services:** Classify content using rule-based or ML models.
- **Entity Extraction & Relationship Mapping Service:** Builds a knowledge graph from textual data.
- **Q&A Service:** Orchestrates search, context, and LLM-based answer generation.

Each module exposes its functionality via a RESTful API (or gRPC) and communicates asynchronously via messaging systems (Kafka or RabbitMQ) to support both real-time and batch processing.

4 Hierarchical Document Model for Scraped Data

4.1 Rationalization

Documents in PDF, Word, or HTML formats typically have nested structures (chapters, sections, paragraphs). Storing them in a **hierarchical** model ensures:

- Preservation of context boundaries for more accurate embeddings.
- Capability to attach metadata (e.g., sentiment, readability) at each level.
- Flexibility to insert images, tables, or subsections without breaking the ingestion pipeline.

4.2 Core Classes and Code Snippets

Document Class Example:

```
1 class Document(DataElement):
2     def __init__(self, title: str, author: Optional[str] = None,
3                 created_date: Optional[datetime] = None,
4                 language: Optional[str] = "English",
5                 translator: Optional[str] = "Original",
6                 description: Optional[str] = None):
7         super().__init__()
8         self.title = title
9         self.author = author
10        self.created_date = created_date or datetime.now()
11        self.language = language
12        self.translator = translator
13        self.description = description
14        self.chapters = []
15        self.cache = {} # For storing computed metadata
16
17    def add_chapter(self, chapter: Chapter):
18        self.chapters.append(chapter)
19
20    def add_section(self, section: Section):
21        if not self.chapters:
22            default_chapter = Chapter(title=self.title, sections=[],
23                                     number=1, author=self.author)
24            self.chapters.append(default_chapter)
25            self.chapters[-1].sections.append(section)
26
27    def aggregate_sentiment(self):
28        if 'sentiment' in self.cache:
29            return self.cache['sentiment']
30        sentiments = [ch.aggregate_sentiment() for ch in self.chapters]
31        self.cache['sentiment'] = (sum(sentiments)/len(sentiments)
32                                   if sentiments else 0)
33        return self.cache['sentiment']
```

Listing 1: Document Class for Storing Hierarchical Structures

Section Class Example:

```
1 class Section(DataElement, DocumentComponent):
2     def __init__(self, title: str, content: List[Union[Paragraph, Image, Table
3         ]],
4         level: int = 1, author: Optional[str] = None):
5         super().__init__()
6         self.title = title
7         self.content = content
8         self.level = level
9         self.subsections = []
10        self.author = author
11        self.cache = {}
12
13    def add_subsection(self, subsection):
14        self.subsections.append(subsection)
15
16    def aggregate_sentiment(self):
17        if 'sentiment' in self.cache:
18            return self.cache['sentiment']
19        sentiments = [p.analyze_sentiment() for p in self.content
20            if isinstance(p, Paragraph)]
21        self.cache['sentiment'] = (sum(sentiments)/len(sentiments)
22            if sentiments else 0)
23        return self.cache['sentiment']
```

Listing 2: Section Class for Nesting Content (Paragraph, Image, Table)

Paragraph Class Example:

```
1 class Paragraph(DataElement):
2     def __init__(self, text: str, style: str = "Normal"):
3         super().__init__()
4         self.id = generate_unique_id()
5         self.text = text
6         self.style = style
7         self.cache = {}
8
9     def analyze_sentiment(self):
10        if 'sentiment' in self.cache:
11            return self.cache['sentiment']
12        blob = TextBlob(self.text)
13        sentiment = blob.sentiment.polarity
14        self.cache['sentiment'] = sentiment
15        return sentiment
16
17    def extract_entities(self):
18        if 'entities' in self.cache:
19            return self.cache['entities']
20        doc = nlp(self.text)
21        entities = [(ent.text, ent.label_) for ent in doc.ents]
22        self.cache['entities'] = entities
23        return entities
```

Listing 3: Paragraph Class with NLP Methods (Sentiment, Entities, etc.)

Tables and images can similarly be modeled with specialized classes.

5 Domain Schema and JSON-Based Modeling

5.1 Rationalization

Real domains have *conceptual knowledge* such as **entities** (e.g., “Income Tax Authority”), **actions** (e.g., “file,” “pay”), and **interactions** (e.g., “check eligibility”). Defining a **Domain**

Schema in JSON/YAML allows:

- Domain experts to define the *who*, *what*, *how* and the conditions.
- The system to parse these structures to guide entity linking and advanced analytics.
- Automated configuration of RASA with domain-specific intents and entities.

5.2 Schema Example

```
1 {
2   "entities": [
3     "individuals",
4     "Income_Tax_Authority",
5     "non_profit_organizations",
6     "donors"
7   ],
8   "actions": [
9     "calculate",
10    "pay",
11    "notify",
12    "file",
13    "maintain",
14    "raise",
15    "request"
16  ],
17  "interactions": [
18    {
19      "interaction": "income_tax",
20      "operands": ["Income_Tax_Authority", "individuals"],
21      "operation_parameters": [
22        {
23          "tasks": "file",
24          "configuration": ["tax forms", "submission deadlines"],
25          "related_verbs": ["filed", "submitted"],
26          "metrics": ["submission date", "form type"],
27          "input": ["completed_income_tax_form", "submission_date"],
28          "output": "Confirms successful submission."
29        },
30        {
31          "tasks": "check_eligibility",
32          "configuration": ["annual income threshold"],
33          "related_verbs": ["eligible", "not_eligible"],
34          "metrics": ["income threshold", "eligibility status"]
35        }
36      ],
37      "conditions": []
38    }
39  ]
40 }
```

Listing 4: Sample JSON Schema for Domain Modeling

5.2.1 Parsing and Validation

```
1 import json
2
3 class DomainSchema:
4     def __init__(self, entities, actions, interactions):
5         self.entities = entities
6         self.actions = actions
7         self.interactions = interactions
8
```



```
9      @classmethod
10      def from_json(cls, json_file):
11          with open(json_file, 'r') as f:
12              data = json.load(f)
13              if "entities" not in data or "actions" not in data or "interactions"
14                  not in data:
15                      raise ValueError("Domain schema missing required keys.")
16          return cls(data["entities"], data["actions"], data["interactions"])
```

Listing 5: Pseudocode for Loading/Validating Domain Schema

6 Modular NLP Framework with Pluggable Components

In many enterprise use cases, simpler rule-based methods are sufficient for well-defined fields, while advanced ML or LLM-based methods are needed for ambiguous or complex semantics. Table 1 provides an overview.

Key Insight: Rule-based methods are efficient for well-defined tasks but become brittle with increased complexity. Many enterprise systems benefit from a hybrid approach—employing rules for predictable patterns and ML models for handling ambiguous cases.

6.1 Document Ingestion & Parsing

Documents are ingested via APIs, email connectors, or batch uploads. Parsing modules (e.g., PDF parsers or OCR systems) convert files into plain text with associated metadata, making it simple to swap components as needed.

6.2 NLP Modules

Once text is extracted, a series of NLP tasks are applied:

- **Classification:** Tagging documents for routing.
- **Entity Recognition:** Extracting names, dates, and other identifiers.
- **Summarization:** Generating abstracts or extracting key points.
- **Other Analyses:** Language detection, sentiment analysis, etc.

A central orchestration layer ensures each module processes documents sequentially and consistently.

6.3 Search Indexing & Storage

Processed documents, along with their annotations, are stored in searchable databases or knowledge graphs. This layer can initially use systems like Elasticsearch and later migrate to vector-based semantic search solutions.

6.4 Query Interface & Applications

Consumer-facing interfaces, such as search UIs or APIs, interact with the processed data. They leverage stored summaries and can even perform on-demand LLM-based Q&A.

Table 1: Comparison of non-ML (rule-based) vs. ML approaches for various text processing applications.

Application	Feature or Use Case	Non-ML (Rule-based)	ML or LLM-based	When ML is Needed
Document Classification	Label document types (invoice, contract, etc.)	Rule-based tagging with keywords or metadata	SVM, Random Forest, BERT	Format variety or subtle domain context
Named Entity Extraction	Identify people, organizations, amounts	Dictionary or regex-based extraction	CRF, BiLSTM-CRF, Transformer-based (spaCy, HuggingFace)	Complex or evolving domain
Relation Extraction	Identify relationships among entities	Rule-based pattern matching	Neural relation classifiers	Contextual or layered relationships
Sentiment Analysis	Polarity/emotion analysis	Lexicon-based scoring with heuristics	Supervised classifiers, few-shot LLM prompting	Sarcasm, domain-specific nuance
Summarization	Condensed document representation	Extractive (TextRank, LexRank), template-based summaries	Abstractive (BART, T5, GPT models)	Complex multi-topic or unstructured documents needing compression
Search / IR	Retrieve relevant documents	TF-IDF, BM25, keyword matching, regex filters	Dense vector search, RAG (retrieval-augmented generation), transformer-based retrieval	Semantic synonyms, large corpora
Q&A Chatbot	Answer user queries	Finite-state scripts, decision trees, FAQ lookups	LLM-based QA, retrieval-augmented generation	Open-ended, ambiguous queries requiring semantic understanding
Grammar Checking	Fix language errors	Hard-coded rules, dictionary lookups, edit-distance corrections	Transformer-based GEC models (e.g., GPT, BERT)	Complex grammatical or contextual errors
Machine Translation	Translate text	Rule-based, phrase-based translation systems	Neural Machine Translation (Transformer NMT)	Idiomatic, domain-specific expressions, complex language contexts
Topic Modeling	Discover latent topics in large corpora	Manual keyword clustering, rule-based taxonomies	Unsupervised ML (LDA, BERTopic, embedding clustering)	Exploring evolving, large datasets for unknown themes

7 Linking Document Content and Domain Schema

7.1 Motivation

- **NER** identifies references to domain **entities** in paragraphs.
- **Action/Verb Extraction** maps text mentions to domain **actions**.
- **Graph Storage** enables advanced queries (e.g., “Which paragraphs mention `Income_Tax_Authority`?”).

7.2 Process Flow

1. **Scrape/Parse Document** → Paragraph objects.
2. **Extract Entities/Verbs** with NLP → match them to domain schema.
3. **Graph Edges** → e.g., `(:Paragraph)-[:MENTIONS]->(:DomainEntity)`.
4. **Vector DB** → store paragraph embeddings.

7.3 Example Linking Code

```
1 def link_paragraph_to_domain(paragraph: Paragraph, domain_schema: DomainSchema,
2   graph):
3     entities_found = paragraph.extract_entities() # e.g. [("Income Tax
4       Authority", "ORG")]
5     paragraph_id = paragraph.id
6
7     for (text, label) in entities_found:
8       matched_entity = fuzzy_match_entity(text, domain_schema.entities)
9       if matched_entity:
10         graph.merge_entity_node(matched_entity)
11         graph.link_paragraph_entity(paragraph_id, matched_entity)
```

Listing 6: Entity Linking Pseudocode

8 Embeddings and Retrieval-Augmented Generation (RAG)

8.1 Paragraph-Level vs. Domain Concept Embeddings

- **Paragraph Embeddings:** Represent chunks of text (typically 2-3 sentences) for detailed semantic search.
- **Domain Concept Embeddings:** Represent domain-specific terms (e.g., “Income Tax Authority”) to bridge queries.

8.2 Implementation Steps

1. **Chunk and Clean:** Split long paragraphs if necessary.
2. **Embed with Transformers:** Use models like *SentenceTransformers* to convert text into vectors.
3. **Store in Vector Database:** Use systems like FAISS or Pinecone for efficient retrieval.

8.3 RAG Query Pipeline

1. **User Query:** Input query, e.g., “How do I file an income tax return?”
2. **Convert Query to Embedding:** Use the same model to embed the query.
3. **ANN Search:** Retrieve top-K similar paragraphs.
4. **Graph Query (Optional):** Perform a knowledge graph traversal if needed.
5. **LLM Prompting:** Fuse the retrieved text with the query for a grounded answer.

8.4 Embeddings and RAG Diagram




Figure 2: Embeddings and RAG Pipeline: From User Query to Grounded LLM Response

9 Detailed “How” Steps & Best Practices

9.1 Ingestion and Parsing

- Create a `Document` object from each source file. If no chapters exist, attach a default `Chapter`.
- Split content into `Section`, `Paragraph`, `Table`, and `Image`.
- Cache NLP outputs (sentiment, entities, keywords) for reuse.

9.2 Domain Schema Loading

- Validate domain schema using `jsonschema` or custom code.
- Maintain a single JSON/YAML file for ease of updates.

9.3 Entity & Action Linking

- Use NER and fuzzy matching to map recognized entities to domain definitions.
- Extract verbs/actions and link text mentions to domain `actions`.

9.4 Embedding Storage

- Process paragraphs in batches for embedding.
- Store resulting vectors in a dedicated vector database.

9.5 RAG Query Processing

- Assemble a prompt by combining top-K retrieved paragraphs with the user’s query.
- Optionally fuse graph-based context with vector search results.

9.6 High-level Flow Diagram for How Steps

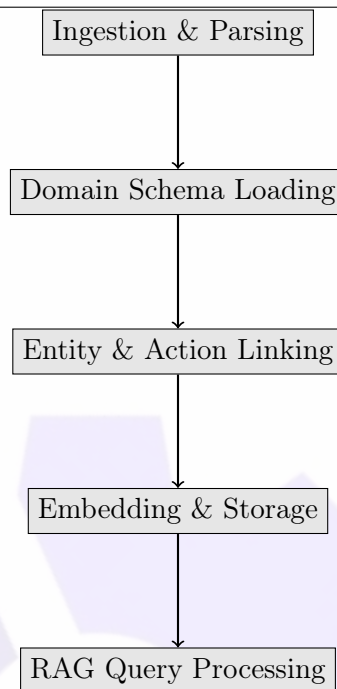


Figure 3: High-level Flow for Ingestion, Domain Processing, and RAG Query Steps

10 LLM Integration with Adapter Pattern

LLMs are integrated through a dedicated Adapter Pattern that abstracts the underlying model. This ensures that the enterprise can choose the optimal model for its needs while maintaining a consistent interface.

10.1 Adapter Pattern Design

- A unified API, e.g. `LLMService.generateText(prompt)`, is provided for all LLM-powered services.
- The adapter is configured at deployment (static model selection) to use a default model (Llama 3.3) or another enterprise-approved model.
- The adapter handles prompt formatting, context injection, and response parsing.
- New adapters can be implemented without altering the core microservice logic, making the integration pluggable.

10.2 Code Snippet: LLM Adapter Interface (Pseudocode)

```
1 class LLMAdapterInterface:
2     def generate_text(self, prompt: str, context: dict = None) -> str:
3         raise NotImplementedError("Subclasses must implement this method")
4
5 class LlamaAdapter(LLMAdapterInterface):
6     def __init__(self, model_path: str):
7         # Load Llama 3.3 model from local disk or endpoint
8         self.model = load_llama_model(model_path)
9
10    def generate_text(self, prompt: str, context: dict = None) -> str:
11        # Format the prompt as needed and call the model
12        formatted_prompt = self.format_prompt(prompt, context)
13        response = self.model.generate(formatted_prompt)
14        return response
```

```

15
16 def load_llama_model(model_path: str):
17     # Pseudocode to load the Llama model
18     pass

```

Listing 7: LLM Adapter Interface Example

11 Logical Components

11.1 Overview

The system comprises multiple logical components including the data ingestion pipeline, graph and vector databases, RASA integration, and LLM inference. These components work in synergy to deliver both structured API calls and open-ended, context-aware responses.

11.2 Service Interaction Diagram

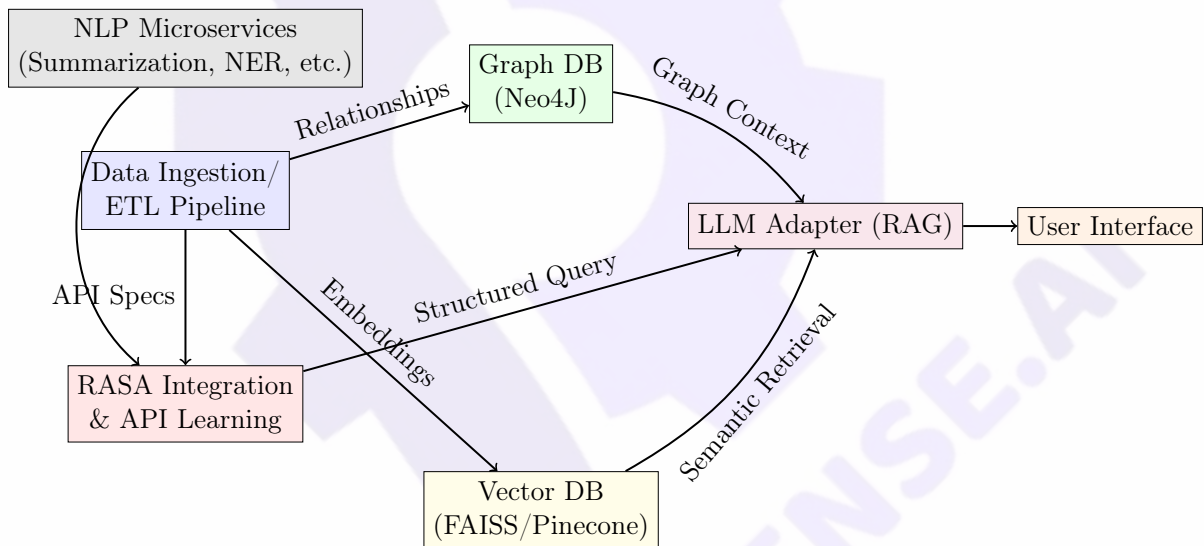


Figure 4: Service Interaction Diagram

12 Real-Time and Batch Processing

12.1 Real-Time Processing

The system leverages an event-driven architecture for real-time processing:

- **Messaging Backbone:** Apache Kafka or RabbitMQ is used to distribute events between microservices.
- **Use Case Example:** Social media or live chat moderation uses real-time streams to trigger sentiment analysis, spam detection, and intent recognition.
- **Scalability:** Each microservice can horizontally scale by consuming messages from partitioned topics.

12.2 Batch Processing

For tasks such as nightly document ingestion or re-indexing:

- **ETL Integration:** The system integrates with an ETL platform to schedule batch jobs.
- **Hybrid Usage:** Batch jobs can read from the same storage as real-time streams, ensuring consistency.
- **Example:** End-of-day processing of purchase orders (POs) where document parsing and indexing are performed in bulk.

13 RASA Integration and Automated API Learning

13.1 Why RASA?

- **Structured Dialogue:** RASA handles intent-based queries, entity extraction, and controls conversation flows for tasks like API calls.
- **Privacy and Customization:** Being open-source, RASA can be self-hosted, ensuring complete data control.
- **LLM Complementarity:** RASA manages deterministic flows, while LLMs handle open-ended questions.

13.2 API Learning: Generating RASA Artifacts

When new APIs (OpenAPI, RAML, WSDL, etc.) are introduced, the system auto-generates:

- RASA Actions (e.g., `action_get_purchase_order`)
- NLU Data (intents and entities based on API specifications)
- External configuration files for credentials and endpoints.

Example: For a GET `/orders/{order_id}` endpoint:

- intent: `get_order_details`
- entity: `order_id`
- action: `action_get_order_details`

13.3 Sample Code for Dynamic Action Generation

```
1 def generate_rasa_actions_from_openapi(openapi_spec_path: str):
2     spec = load_openapi_spec(openapi_spec_path)
3
4     for path, methods in spec["paths"].items():
5         for method, details in methods.items():
6             action_name = f"action_{method.lower()}_{path_to_identifier(path)}"
7             action_code = f"""
8 class {action_name.title().replace('_', '')}(Action):
9     def name(self):
10         return "{action_name}"
11
12     def run(self, dispatcher, tracker, domain):
13         # Extract slots and call the API
14         return []
15 """
16         write_action_code_to_file(action_code)
```

Listing 8: Pseudocode for Generating RASA Actions from an OpenAPI Spec

13.4 Connecting Domain Schema to RASA

- Map domain schema actions to RASA intents and actions.
- Use domain schema entities as synonyms in RASA NLU configurations.

14 LLM Integration in the Framework

LLMs are integrated as modular, swappable components that can enhance tasks such as summarization, question-answering, and classification. Key aspects include:

- **Service Interface:** A unified API (e.g., `LLMService.summarize(text)`) abstracts the underlying model, enabling dynamic switching between proprietary (GPT-4, PaLM) and open-source models (LLaMA, GPT-J).
- **Fine-Tuning:** The framework allows for continuous improvement by fine-tuning models on enterprise-specific data.
- **Integration Points:** LLMs can be employed in summarization, Q&A, classification, and extraction tasks, with a consistent input/output format that fits into the overall pipeline.

15 Additional Considerations

Building and maintaining an enterprise NLP system requires attention to several operational best practices:

15.1 Evaluation and Benchmarking

- **Metrics:** Use task-appropriate metrics (accuracy, F1, ROUGE, NDCG) on representative datasets.
- **Test Cases:** Develop enterprise-specific golden datasets to capture edge cases.
- **Automated Regression:** Incorporate continuous testing to detect performance drops.
- **Classification Tasks:** Use accuracy, precision, recall, and F1-score.
- **Generative Tasks:** Use BLEU, ROUGE, and NDCG for summarization and translation.
- **Search Relevance:** Measure precision@K and NDCG.

15.2 Continuous Learning and Adaptation

- **Active Learning:** When uncertainty is high, the system flags inputs for human annotation to improve training datasets.
- **Feedback Loops:** Use evaluation results to refine models periodically.
- **Scheduled Retraining:** Maintain regular retraining cycles with updated enterprise data.
- **User Feedback Collection:** End-user actions (clicks, ratings, corrections) are collected to gauge quality.
- **Model Drift Detection:** Monitor production metrics and compare against validation benchmarks.
- **Automated Retraining:** ETL pipelines aggregate feedback data to trigger periodic fine-tuning of models.

15.3 Explainability and Human-in-the-Loop

- **XAI:** Implement methods (e.g., feature importance, attention visualization) to explain model decisions.
- **Human Oversight:** Provide interfaces for human verification and correction of low-confidence predictions.
- **Auditability:** Maintain detailed logs to trace and review system decisions.

16 Deployment, Scalability, and Testing

16.1 Error Handling and Monitoring

- Log schema validation errors, parsing failures, and NER mismatches.
- Implement auto-retry mechanisms for critical failures.

16.2 Scaling the Graph and Vector DB

- Use graph partitioning (e.g., Neo4J clustering) and vector index sharding for high-scale environments.

16.3 Scaling LLM Inference

- Distribute inference endpoints and manage token window limits via summarization or chunking.

16.4 Testing and Validation

- **Unit Tests:** Validate document structures, domain schema parsing, and entity extraction.
- **Integration Tests:** Verify document-to-embedding-to-RAG query pipelines.
- **Load Testing:** Test concurrency on vector retrieval, graph queries, and LLM inference.

17 Security and Access Control

- Implement Role-Based Access Control (RBAC) for schema modifications and pipeline executions.
- Ensure transport security (TLS/SSL) for all data transmissions.
- **Security:** Encrypt sensitive data at rest and enforce fine-grained permissions. Ensure encryption, access control, and data privacy, especially when **integrating external LLM APIs**.
- Maintain audit logs for changes in the domain schema, knowledge graph, and user queries.
- **Performance Optimization:** Monitor latency and use caching or batching as appropriate.

18 Conclusion, Critical Gaps, and Future Directions

This Unified Enterprise Knowledge Hub provides a robust framework integrating a hierarchical document model, domain schema, knowledge graph, embedding-based retrieval, and the synergy of RASA and LLM for both structured and open-ended queries. Key benefits include:

- *Structured + Unstructured Synergy:* Blends paragraph-level text with domain-level concepts.

-
- *LLM Hallucination Mitigation:* Grounds LLM responses in retrieved document passages.
 - *Scalable Architecture:* Modular design with optional graph, vector, and caching layers.
 - *Faster API Integration:* Automated RASA artifact generation from API specifications.

18.1 Critical Gaps and Recommendations

- **Domain Modeling Consistency:** Adopt a centralized, versioned schema registry.
- **Temporal and Causal Inference:** Consider integrating time-stamped node states and Bayesian/GNN modules.
- **LLM Hallucination:** Provide source citation features and post-check responses against domain constraints.
- **Workflow Orchestration:** Expand the schema to define multi-step processes with a dedicated workflow engine.
- **Performance Optimization:** Use caching, summarization, or model quantization to manage resource usage.

18.2 Future Directions

- **Temporal/Causal Reasoning:** Integrate temporal graphs for domains requiring event tracking.
- **Simulation and Workflow:** Model and orchestrate multi-step business processes.
- **LLM Fine-Tuning:** Further adapt LLMs to specialized enterprise jargon.
- **Multi-Modal Integration:** Incorporate speech, video, and image processing to broaden data sources.

19 Final Reflection

This document outlines a comprehensive, scalable architecture for an enterprise knowledge hub that bridges raw data and actionable intelligence. With enhancements in document modeling, domain schema standardization, graph and vector integration, and RASA-LLM synergy, the system is poised for production-scale deployment and future innovation.